

Assignment #3

Modulation Classification

Name1: Moustafa Ibrahim Tohamy
ID1: 3851

Name2: Zeyad Ezzat
ID2: 4492

Introduction:

A synthetic dataset, generated with GNU Radio, consisting of 10 modulations. This is a variable-SNR dataset with moderate LO drift, light fading, and numerous different labeled SNR increments for use in measuring performance across different signal and noise power scenarios.

Steps:

Data has been extracted using pickle.

The Code for dataset extraction:

```
import pickle
import numpy as np

'''
    the data set has 10 modulations
    for each modulation we have 20 snr
    for each modulation/snr we have 1000 instance
    *****
        for example
    8 PSK modulation has 20 snr values from -20 to 18 with step value = 2
    for each value of the snr for 8 PSK modulation we have 1000 example
    and the same for each modulation
    so for every modulation we have 20(snr)*1000=20000 example
    *****
    the dataset = 10 (modulation)* 20000(example for each modulation)=200000 example
'''

a=open("/content/gdrive/My Drive/a/RML2016.10b.dat",'rb')
u = pickle._Unpickler(a)
u.encoding = 'latin1'
Xd = u.load()

# Xd = pickle.load(a)
snrs,mods = map(lambda j: sorted(list(set(map(lambda x: x[j], Xd.keys())))), [1,0])
X = []
lbl = []

j=0
```

```

classes_index=0
mod_to_int={}
for mod in mods:
    j+=1
    print('mod = ',mod)
    mod_to_int[j] = mod

    for snr in snrs:
        """
        classes are the labels dataset each modulation will have a number from 1 to 11
        mod to int is to map between numbers and classes
        dataset_snrs is all snr as a one vector to be more easier to deal with

        """
        classes_index+=1000
        print(Xd[mod,snr].shape)
        X.append(Xd[(mod,snr)])
        for i in range(Xd[(mod,snr)].shape[0]): lbl.append((mod,snr))
X = np.vstack(X)
import matplotlib.pyplot as plt
y=np.zeros(shape=128)
for i in range(0,128):
    y[i]=i
plt.plot(y,X[-8000][0])
plt.plot(y,X[-8000][1])
plt.show()

```

Splitting Data

```

np.random.seed(2016)
n_examples = X.shape[0]
n_train = n_examples * 0.5
train_idx = np.random.choice(range(0,n_examples), size= int(n_train), replace=False)
test_idx = list(set(range(0,n_examples))-set(train_idx))
X_train = X[train_idx]
X_test = X[test_idx]
# one hot encoding for multiclass classification since there are 11 classes (11 modulation techniques)
def to_onehot(yy):
    yy1 = np.zeros([len(yy), max(yy)+1])
    yy1[np.arange(len(yy)),yy] = 1
    return yy1
Y_train = to_onehot(list(map(lambda x: mods.index(lbl[x][0]), train_idx)))
Y_test = to_onehot(list(map(lambda x: mods.index(lbl[x][0]), test_idx)))
in_shp = list(X_train.shape[1:])
# print(X_train.shape, " ", in_shp)
print(in_shp+[1])
classes = mods

```

***Note: The raw features were used in our training model as the derivatives provided so low accuracy.**

```
#Defining some variables (Epochs ,Dropout rate , batch size)
dr = 0.5
# Set up some params for training
nb_epoch = 100 # number of epochs to train on
batch_size = 1024 # training batch size
```

1st - Fully connected neural network

```
#Build the N-Network
model = keras.models.Sequential()
model.add(Reshape(in_shp+[1], input_shape=in_shp))
model.add(Dropout(dr))
model.add(Flatten())
model.add(Dense(128, activation='relu', kernel_initializer='he_normal', name="dense1"))
model.add(Dropout(dr))
model.add(Dense(128, activation='relu', kernel_initializer='he_normal', name="dense2"))
model.add(Dropout(dr))
model.add(Dense(64, activation='relu', kernel_initializer='he_normal', name="dense3"))
model.add(Dropout(dr))
model.add(Dense(64, activation='relu', kernel_initializer='he_normal', name="dense4"))
model.add(Dropout(dr))
model.add(Dense(len(classes), kernel_initializer='he_normal', name="dense5"))
model.add(Activation('softmax'))
#Reshaping layer to output one of our 10 classes
model.add(Reshape([len(classes)]))

#Compile
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
model.summary()
```

Perform training ...

```
filepath = 'convmodrecnets_CNN2_0.5.wts.h5'
history=model.fit(X_train,
    Y_train,
    batch_size=batch_size,
    epochs=nb_epoch,
    verbose=2,
    validation_split=0.05,
    callbacks = [
        keras.callbacks.ModelCheckpoint(filepath, monitor='val_loss', verbose=0,
save_best_only=True, mode='auto'),
        keras.callbacks.EarlyStopping(monitor='val_loss', patience=5, verbose=0, mode='auto')
    ])
# we re-load the best weights once training is finished
model.load_weights(filepath)
```


```
score = model.evaluate(X_test, Y_test, batch_size=batch_size)
print(model.metrics_names)
print(score)
```




OUTPUT:



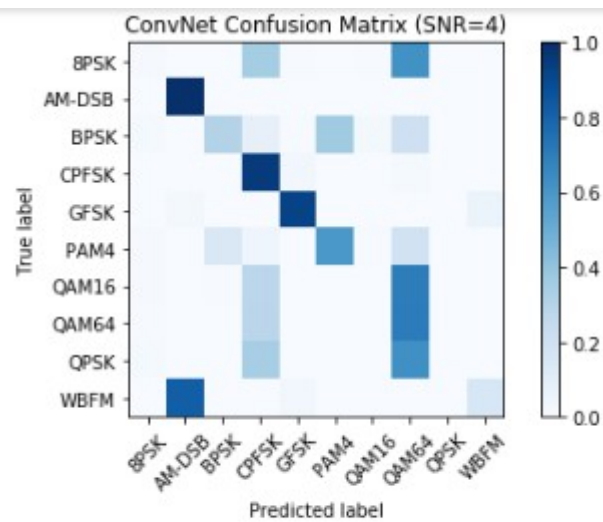
```
['loss', 'acc'] [1.2756312332344055, 0.46331166670481366]
```



Overall Accuracy for each corresponding SNR value:

Overall Accuracy:	0.1051031995190702
Overall Accuracy:	0.10674194841641053
Overall Accuracy:	0.11280701169727064
Overall Accuracy:	0.12269363426313498
Overall Accuracy:	0.14665242640678192
Overall Accuracy:	0.21109963689663214
Overall Accuracy:	0.27046759366807155
Overall Accuracy:	0.3509035641513846
Overall Accuracy:	0.478809888718598
Overall Accuracy:	0.5091116173120729
Overall Accuracy:	0.46899654211568803
Overall Accuracy:	0.4675099866844208
Overall Accuracy:	0.463636061321539
Overall Accuracy:	0.46979287489643745
Overall Accuracy:	0.46533010256580476
Overall Accuracy:	0.4671100764881942
Overall Accuracy:	0.46211969908794354
Overall Accuracy:	0.46574888041134516
Overall Accuracy:	0.4626213430296561
Overall Accuracy:	0.4621969570305969









CNN

CNN using PDF Architecture

```

model2 = keras.models.Sequential()
model2.add(Reshape(in_shp+[1], input_shape=in_shp))
# no padding for the height and add padding to width (2 more columns)
model2.add(ZeroPadding2D((0, 2)))
model2.add(Conv2D(64, (1, 3), padding='valid', activation='relu', name="conv1",
                  kernel_initializer='glorot_uniform', data_format="channels_last"))
# Adding dropout to inputs to next layer to avoid over fitting
model2.add(Dropout(dr))
model2.add(ZeroPadding2D((0, 2)))

```

```

model2.add(Conv2D(16, (2, 3), padding='valid', activation='relu', name="conv2",
kernel_initializer='glorot_uniform', data_format="channels_last"))
model2.add(Dropout(dr))
# The coming layer is dense so we need to flatten our inputs
model2.add(Flatten())
model2.add(Dense(128, activation='relu', kernel_initializer='he_normal', name="dense1"))
model2.add(Dropout(dr))
model2.add(Dense(len(classes), kernel_initializer='he_normal', name="dense2" ))
model2.add(Activation('softmax'))
model2.add(Reshape([len(classes)]))
model2.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
model2.summary()

filepath = 'convmodrecnets_CNN2_0.5.wts.h5'
history=model2.fit(X_train,
Y_train,
batch_size=batch_size,
epochs=nb_epoch,
verbose=2,
validation_split=0.05,
callbacks = [
keras.callbacks.ModelCheckpoint(filepath, monitor='val_loss', verbose=0,
save_best_only=True, mode='auto'),
keras.callbacks.EarlyStopping(monitor='val_loss', patience=10, verbose=0, mode='auto')
])
# we re-load the best weights once training is finished
model2.load_weights(filepath)

```


OutPut:

Last Epoch result:

Epoch 54/100 - 13s - loss: 1.2126 - acc: 0.4863 - val_loss: 1.1303 -
val_acc: 0.5225

[1.0480517408156396, 0.5663266666603088]

Overall Accuracy: 0.10436844566161245
Overall Accuracy: 0.10640785781103836
Overall Accuracy: 0.11457326623787782
Overall Accuracy: 0.1266115259685899
Overall Accuracy: 0.145384153260797
Overall Accuracy: 0.22595689396715413
Overall Accuracy: 0.39325656257259484
Overall Accuracy: 0.5348899355312824
Overall Accuracy: 0.625708002931965
Overall Accuracy: 0.6798204475412033
Overall Accuracy: 0.7117198777990399
Overall Accuracy: 0.7327896138482024
Overall Accuracy: 0.7364570516444415
Overall Accuracy: 0.7457166528583264
Overall Accuracy: 0.7481992896737146
Overall Accuracy: 0.7464582640505487
Overall Accuracy: 0.7422275480993276
Overall Accuracy: 0.7449991706750705
Overall Accuracy: 0.7480068052173333
Overall Accuracy: 0.7474001003176727



Then we tried to modify the architecture and parameters as :

#Defining some variables (Epochs ,Dropout rate , batch size)

dr = 0.05

Set up some params for training

nb_epoch = 100 # number of epochs to train on

batch_size = 1024 # training batch size

Model Code:

```
model3 = keras.models.Sequential()
model3.add(Reshape(in_shp+[1], input_shape=in_shp))
# no padding for the height and add padding to width (2 more columns)
model3.add(ZeroPadding2D((0, 2)))
model3.add(Conv2D(64, (1, 3), padding='valid', activation='relu', name="conv1",
                  kernel_initializer='glorot_uniform', data_format="channels_last"))
# Adding dropout to inputs to next layer to avoid over fitting
model3.add(Dropout(dr))
model3.add(ZeroPadding2D((0, 2)))
model3.add(Conv2D(64, (2, 3), padding='valid', activation='relu', name="conv2",
                  kernel_initializer='glorot_uniform', data_format="channels_last"))
model3.add(Dropout(dr))
# The coming layer is dense so we need to flatten our inputs
model3.add(Flatten())
model3.add(Dense(128, activation='relu', kernel_initializer='he_normal', name="dense1"))
model3.add(Dropout(dr))
model3.add(Dense(len(classes), kernel_initializer='he_normal', name="dense2"))
```

```
model3.add(Activation('softmax'))
model3.add(Reshape([len(classes)]))
model3.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

Train on 570000 samples, validate on 30000 samples

Epoch 1/100 - 14s - loss: 1.9088 - acc: 0.2465 - val_loss: 1.6231 - val_acc: 0.3500
Epoch 2/100 - 14s - loss: 1.4913 - acc: 0.3967 - val_loss: 1.3809 - val_acc: 0.4438
Epoch 3/100 - 14s - loss: 1.3094 - acc: 0.4631 - val_loss: 1.2695 - val_acc: 0.4779
Epoch 4/100 - 14s - loss: 1.2324 - acc: 0.4863 - val_loss: 1.2272 - val_acc: 0.4869
Epoch 5/100 - 14s - loss: 1.2044 - acc: 0.4940 - val_loss: 1.2255 - val_acc: 0.4787
Epoch 6/100 - 13s - loss: 1.1844 - acc: 0.5012 - val_loss: 1.1840 - val_acc: 0.4957
Epoch 7/100 - 13s - loss: 1.1689 - acc: 0.5066 - val_loss: 1.1785 - val_acc: 0.5038
Epoch 8/100 - 13s - loss: 1.1526 - acc: 0.5134 - val_loss: 1.1567 - val_acc: 0.5145
Epoch 9/100 - 13s - loss: 1.1376 - acc: 0.5187 - val_loss: 1.1464 - val_acc: 0.5081
Epoch 10/100 - 13s - loss: 1.1265 - acc: 0.5235 - val_loss: 1.1360 - val_acc: 0.5149
Epoch 11/100 - 14s - loss: 1.1184 - acc: 0.5276 - val_loss: 1.1448 - val_acc: 0.5076
Epoch 12/100 - 13s - loss: 1.1126 - acc: 0.5304 - val_loss: 1.1380 - val_acc: 0.5201
Epoch 13/100 - 13s - loss: 1.1053 - acc: 0.5345 - val_loss: 1.1332 - val_acc: 0.5269
Epoch 14/100 - 13s - loss: 1.0980 - acc: 0.5378 - val_loss: 1.1163 - val_acc: 0.5309
Epoch 15/100 - 13s - loss: 1.0893 - acc: 0.5426 - val_loss: 1.1165 - val_acc: 0.5341
Epoch 16/100 - 13s - loss: 1.0774 - acc: 0.5500 - val_loss: 1.1073 - val_acc: 0.5391
Epoch 17/100 - 14s - loss: 1.0676 - acc: 0.5550 - val_loss: 1.0947 - val_acc: 0.5433
Epoch 18/100 - 13s - loss: 1.0606 - acc: 0.5586 - val_loss: 1.1009 - val_acc: 0.5376
Epoch 19/100 - 13s - loss: 1.0538 - acc: 0.5625 - val_loss: 1.0915 - val_acc: 0.5471
Epoch 20/100 - 14s - loss: 1.0457 - acc: 0.5658 - val_loss: 1.0899 - val_acc: 0.5465
Epoch 21/100 - 13s - loss: 1.0400 - acc: 0.5688 - val_loss: 1.0796 - val_acc: 0.5489
Epoch 22/100 - 13s - loss: 1.0356 - acc: 0.5709 - val_loss: 1.1050 - val_acc: 0.5414
Epoch 23/100 - 14s - loss: 1.0311 - acc: 0.5728 - val_loss: 1.0797 - val_acc: 0.5503
Epoch 24/100 - 13s - loss: 1.0275 - acc: 0.5740 - val_loss: 1.0836 - val_acc: 0.5480
Epoch 25/100 - 13s - loss: 1.0220 - acc: 0.5772 - val_loss: 1.0808 - val_acc: 0.5509
Epoch 26/100 - 13s - loss: 1.0173 - acc: 0.5792 - val_loss: 1.0796 - val_acc: 0.5504
Epoch 27/100 - 13s - loss: 1.0120 - acc: 0.5813 - val_loss: 1.0771 -



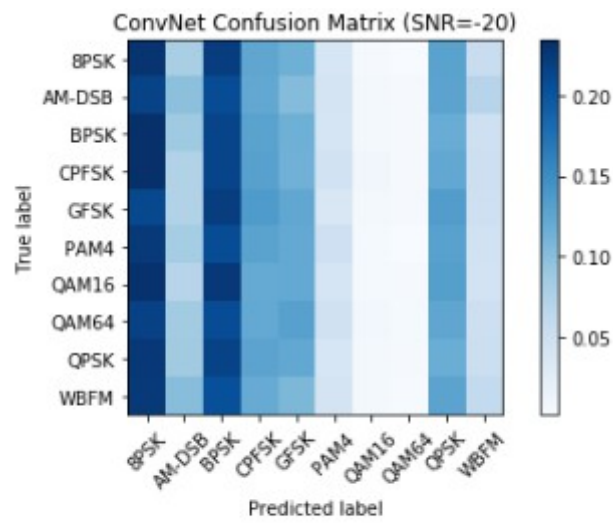
val_acc: 0.5535
Epoch 28/100 - 13s - loss: 1.0066 - acc: 0.5842 - val_loss: 1.0786 -
val_acc: 0.5541
Epoch 29/100 - 14s - loss: 1.0000 - acc: 0.5871 - val_loss: 1.0800 -
val_acc: 0.5526




Epoch 30/100 - 13s - loss: 0.9911 - acc: 0.5909 - val_loss: 1.0702 -
val_acc: 0.5540
Epoch 31/100 - 13s - loss: 0.9837 - acc: 0.5944 - val_loss: 1.0645 -
val_acc: 0.5594
Epoch 32/100 - 13s - loss: 0.9739 - acc: 0.5992 - val_loss: 1.0648 -
val_acc: 0.5560
Epoch 33/100 - 13s - loss: 0.9663 - acc: 0.6026 - val_loss: 1.0734 -
val_acc: 0.5593
Epoch 34/100 - 14s - loss: 0.9584 - acc: 0.6063 - val_loss: 1.0580 -
val_acc: 0.5653
Epoch 35/100 - 14s - loss: 0.9512 - acc: 0.6093 - val_loss: 1.0583 -
val_acc: 0.5652
Epoch 36/100 - 13s - loss: 0.9453 - acc: 0.6113 - val_loss: 1.0776 -
val_acc: 0.5547
Epoch 37/100 - 13s - loss: 0.9430 - acc: 0.6124 - val_loss: 1.0658 -
val_acc: 0.5657
Epoch 38/100 - 13s - loss: 0.9369 - acc: 0.6152 - val_loss: 1.0655 -
val_acc: 0.5662
Epoch 39/100 - 13s - loss: 0.9323 - acc: 0.6171 - val_loss: 1.0682 -
val_acc: 0.5604
Epoch 40/100 - 13s - loss: 0.9276 - acc: 0.6198 - val_loss: 1.0702 -
val_acc: 0.5650
Epoch 41/100 - 14s - loss: 0.9241 - acc: 0.6201 - val_loss: 1.0706 -
val_acc: 0.5628
Epoch 42/100 - 14s - loss: 0.9188 - acc: 0.6230 - val_loss: 1.0790 -
val_acc: 0.5608
Epoch 43/100 - 14s - loss: 0.9152 - acc: 0.6251 - val_loss: 1.0799 -
val_acc: 0.5628
Epoch 44/100 - 14s - loss: 0.9104 - acc: 0.6271 - val_loss: 1.0899 -
val_acc: 0.5628

Print SNRS Accuracies

Overall Accuracy: 0.10289893794669695
Overall Accuracy: 0.10590672190298009
Overall Accuracy: 0.11583963741793582
Overall Accuracy: 0.14144593644309011
Overall Accuracy: 0.18433348908617583
Overall Accuracy: 0.27122822212598685
Overall Accuracy: 0.36730494806358477
Overall Accuracy: 0.4687176403781274
Overall Accuracy: 0.596021856466982
Overall Accuracy: 0.7141900040198311
Overall Accuracy: 0.7783596871118273
Overall Accuracy: 0.8067909454061252
Overall Accuracy: 0.8161351468191945
Overall Accuracy: 0.8202485501242751

Graphs:





After trying some more model modifications:

```
model4 = keras.models.Sequential()
model4.add(Reshape(in_shp+[1], input_shape=in_shp))
# no padding for the height and add padding to width (2 more columns)
model4.add(ZeroPadding2D((0, 2)))
model4.add(Conv2D(256, (1, 3), padding='valid', activation='relu', name="conv1",
                 kernel_initializer='glorot_uniform', data_format="channels_last"))
# Adding dropout to inputs to next layer to avoid over fitting
model4.add(Dropout(dr))
model4.add(ZeroPadding2D((0, 2)))
model4.add(Conv2D(256, (2, 3), padding='valid', activation='relu', name="conv2",
                 kernel_initializer='glorot_uniform', data_format="channels_last"))
model4.add(Dropout(dr))
# The coming layer is dense so we need to flatten our inputs
model4.add(Flatten())
model4.add(Dense(256, activation='relu', kernel_initializer='he_normal', name="dense1"))
model4.add(Dropout(dr))
model4.add(Dense(len(classes), kernel_initializer='he_normal', name="dense2" ))
model4.add(Activation('softmax'))
model4.add(Reshape([len(classes)]))
model4.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

2

Last epoch result:

```
Epoch 25/100 - 111s - loss: 0.8556 - acc: 0.6522 - val_loss: 1.1813 -
val_acc: 0.5553
```